

This lab is intended to get the team familiar with creating a pod in the team namespace. There is no diagnosis or problem to be researched. The lab is complete once the pod is created in the team namespace.

Resources

- K8 yaml - [house.yaml](#)
- Dockerfile - [Dockerfile](#)

Useful information

Item	Value
cpu:	50m
memory:	50Mi
image:	ibmicpcoc/house:latest
ports	none
Docker	CMD ["/bin/bash", "-c", "./house.sh"]

Task description

Download the resource K8 yaml file.

Edit and save the file after replacing all references of **<team>** with your team name.

Create the K8 objects using `oc create`

Did the pod deploy successfully? If not, correct the issue and re-create the K8 objects.

To create the pod use the command: **oc create -f <file>**; (replace <file> with the name of the yaml file you have saved an edited.)

Diagnosis

No diagnosis is necessary for this lab. A new pod should be created after editing the yaml file and using the `oc create` command.

Problem discovered

N/A

Resolution

Edit the yaml file and modify all references of `<team>` to your team name.

Example yaml file that needs to be edited.

```
--- # Fast Start :: Problem Diagnosis and Troubleshooting Lab
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: <team>-house
  namespace: <team>
  labels:
    app: <team>-house
```

```
spec:
  selector:
    matchLabels:
      app: <team>-house
  replicas: 1
  template:
    metadata:
      labels:
        app: <team>-house
    spec:
      containers:
      - name: <team>-house
        image: ibmicpcoc/house:latest
        imagePullPolicy: Always
        command: ["/bin/bash", "-c", "/app/avail.sh"]
        env:
          - name: APP_NAMESPACE
            valueFrom:
              fieldRef:
                fieldPath: metadata.namespace
          - name: APP_NAME
            valueFrom:
              fieldRef:
                fieldPath: metadata.name
          - name: COLLECTOR_CONFIG
            valueFrom:
              configMapKeyRef:
                name: <team>-collector-config
                key: COLLECTOR_CONFIG
          - name: INSTRUCTOR_CONFIG
            valueFrom:
              configMapKeyRef:
                name: <team>-collector-config
                key: INSTRUCTOR_CONFIG
      resources:
        requests:
          cpu: 100m
          memory: 100Mi
```

Saved the modified file and create the pod "house".

```
Command to create the K8 objects:
oc create -f house.yaml
```

```
Result output:
deployment.apps/house created
```

Verify the pod deployed successfully.

```
Command to get pods in namespace:
oc -n <team> get pods           # change <team> to your team namespace
```

All references to "team" or <team> should be replaced with your team name which is the same as your namespace.

Use the debug flow to guide the steps you should attempt in diagnosis of the issue.

Resources

- K8 yaml - [baker.yaml](#)
- Dockerfile - [Dockerfile](#)

Useful information

Item	Value
cpu:	50m
memory:	50Mi
image:	ibmicpcoc/baker:latest
ports	none
Docker	CMD ["/bin/bash", "-c", "./baker.sh"]

Task description
Download the resource K8 yaml file.
Edit and save the file after replacing all references of <team> with your team name / namespace.
Research why the pod did not deploy.
Resolve the issue and create the K8 objects.
Did the pod deploy successfully? If not, correct the issue and re-create the K8 objects.

Deployment.spec.template.spec.containers expects an array of entires.

Arrays are defined with a hyphen.

Review and compare the **house.yaml** file for an example of properly defined K8 objects.

Diagnosis

When attempting to create the pod the yaml is not properly defined. This error message is being shown:

error: error validating "baker.yaml": error validating data: ValidationError(Deployment.spec.template.spec.containers): invalid type for io.k8s.api.core.v1.PodSpec.containers: got "map", expected "array"; if you choose to ignore these errors, turn validation off with --validate=false

Problem discovered

The Deployment.spec.template.spec.containers portion of the yaml file is not properly formatted. Got "map", expected "array". Container does not have an array of entires.

Resolution

Edit the yaml file and correct the definition to include a hyphen before the "name:" parameter of the containers portion.

```
Example saved file with hyphen (portion of file shown below)

apiVersion: apps/v1
kind: Deployment
metadata:
  name: <team>-baker
  namespace: pink
  labels:
    app: <team>-baker
spec:
  selector:
    matchLabels:
      app: <team>-baker
  replicas: 1
  template:
    metadata:
      labels:
        app: <team>-baker
    spec:
      containers:
      - name: <team>-baker          <=== Add the hyphen to this line
        image: ibmicpoc/baker:latest
        imagePullPolicy: Always
```

Saved the modified file and create the pod "baker".

```
Command to create the K8 objects:
  oc create -f baker.yaml

Result output:
  deployment.apps/baker created

----

Verify the pod deployed successfully.

Command to get pods in namespace:
  oc -n <team> get pods          # change <team> to your team namespace
```

All references to "team" or <team> should be replaced with your team name which is the same as your namespace.

Use the debug flow to guide the steps you should attempt in diagnosis of the issue.

Useful informationK8 yaml - [carbon.yaml](#)Dockerfile - [Dockerfile](#)

Item	Value
spec.template.spec.containers[*].resources.request.cpu	100m
spec.template.spec.containers[*].resources.request.memory:	100Mi
spec.template.spec.containers[*].image:	ibmicpcoc/carbon:latest
spec.template.spec.containers[*].ports	none
Docker CMD	["/bin/bash", "-c", "./carbon.sh"]

Task description

Within your team namespace diagnose the pod that begins with **<team> -carbon**

Use the label option -l app=<team>-carbon when getting the pod information.

Download the resource K8 yaml file.

Edit and save the file after replacing all references of <team> with your team name / namespace.

Create the K8 objects.

Did the pod deploy successfully? If not, correct the issue and re-create the K8 objects.

- Describe the pod.
- Get events from the namespace, oc get events -n **<team>**
- A single cpu is defined with 1000m. The container cpu resources should use **1/10** of a cpu.
- Editing a running pod is another method to change the pod. Use the command KUBE_EDITOR="nano" oc edit deployment/**<team>-carbon** and edit the running pod. Nano is the editor defined in the command. Remove the KUBE_EDITOR parm to use the default editor on your machine.

Diagnosis

When attempting to deploy the pod the yaml file is not properly defined.

Check the Pod status

Command:

```
oc -n <team> get pods -l app=<team>-carbon. # replace <team>
```

Example output:

NAME	READY	STATUS	RESTARTS	AGE
pink-carbon-5c96bc649-tjnhb	0/1	Pending	0	2m

Describe the pod

```

Name:                pink-carbon-5c96bc649-tjnhb
Namespace:           pink
Priority:              0
PriorityClassName:    <none>
Node:                 <none>
Labels:               app=pink-carbon
                     pod-template-hash=175267205
Annotations:          kubernetes.io/psp=ibm-privileged-psp
Status:               Pending
IP:
Controlled By:        ReplicaSet/pink-carbon-5c96bc649
Containers:
  pink-carbon:
    Image:             ibmicpcoc/carbon:latest
    Port:              <none>
    Host Port:         <none>
    Requests:
      cpu:              25
      memory:           100Mi
    Environment:
      APP_NAMESPACE:     pink (v1:metadata.namespace)
      APP_NAME:          pink-carbon-5c96bc649-tjnhb (v1:metadata.name)
      COLLECTOR_CONFIG:  <set to the key 'COLLECTOR_CONFIG' of config map 'pink-collector-config'> Optional:
false
      INSTRUCTOR_CONFIG: <set to the key 'INSTRUCTOR_CONFIG' of config map 'pink-collector-config'> Optional:
false
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-mq64m (ro)
Conditions:
  Type           Status
  PodScheduled   False
Volumes:
  default-token-mq64m:
    Type:          Secret (a volume populated by a Secret)
    SecretName:     default-token-mq64m
    Optional:       false
QoS Class:        Burstable
Node-Selectors:    <none>
Tolerations:       node.kubernetes.io/memory-pressure:NoSchedule
                   node.kubernetes.io/not-ready:NoExecute for 300s
                   node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type           Reason             Age           From           Message
  ----           -
  Warning        FailedScheduling    58s (x121 over 5m)  default-scheduler  0/4 nodes are available: 4 Insufficient cpu.
$

```

In the "Events" section review the "Message" from the entry with "Type" Warning and "Reason" FailedScheduling

```

... 0/4 nodes are available: 4 Insufficient cpu.
$

```

Example of Get Events in namespace

Command:

```
oc -n <team> get events
```

Example output:

LAST SEEN	FIRST SEEN	COUNT	NAME	KIND	SUBJECT
7m	7m	1	pink-carbon.157belef7ad1a77	Deployment	
Normal	ScalingReplicaSet		deployment-controller	Scaled up replica set pink-carbon-5c96bc649 to 1	
7m	7m	1	pink-carbon-5c96bc649.157belef85494ba	ReplicaSet	
Normal	SuccessfulCreate		replicaset-controller	Created pod: pink-carbon-5c96bc649-tjnhb	
2m	7m	121	pink-carbon-5c96bc649-tjnhb.157belef858b4b3	Pod	
Warning	FailedScheduling		default-scheduler	0/4 nodes are available: 4 Insufficient cpu.	

Problem discovered

Events output indicates the pod is FailedScheduling because there are not enough CPU resources available.

Resolution

At least two methods exist to correct the issue.

The first method is deleting the old pod, edit the yaml file, and re-create the pod.

This approach is later referred to as: delete-create-pod

Edit the yaml file and modify *cpu* to decrease the amount of cpu to 10% of a single CPU.

Delete the running pod

Command to delete the existing pod:

```
oc delete -f carbon.yaml
```

Result output:

```
deployment.apps "carbon" deleted
```

Edit file carbon.yaml (only a portion of file shown below)

```
spec:
  selector:
    matchLabels:
      app: <team>-carbon
  replicas: 1
  template:
    metadata:
      labels:
        app: <team>-carbon
    spec:
      containers:
        - name: <team>-carbon
```

```

    image: ibmicpcoc/carbon:latest
    resources:
      requests:
        cpu: 25000m          <=== change value to 100m
        memory: 100Mi

```

Create the k8 deployment

```

Command:
  oc create -f carbon.yaml

Result output:
deployment.apps/<team>-carbon created

```

The second method is editing the running pod. Edit and save edit the file, and re-create the pod.

This approach is later referred as: edit-running-pod

Edit the running pod. The kubernetes object content is available in the editor (shown below). Note the content has both the spec: and status: sections.

Locate the line cpu: "25" and change the line to cpu: 100m (without quotes)

```

Command to edit the running pod:
  KUBE_EDITOR="nano" oc edit deployment/<team>-carbon      # replace <team>

Content shown when editor is open.  The pink-carbon deployment is being shown:

# Please edit the object below. Lines beginning with a '#' will be ignored,
# and an empty file will abort the edit. If an error occurs while saving this file will be
# reopened with the relevant failures.
#
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  annotations:
    deployment.kubernetes.io/revision: "1"
  creationTimestamp: 2019-01-21T14:01:56Z
  generation: 1
  labels:
    app: pink-carbon
  name: pink-carbon
  namespace: pink
  resourceVersion: "5834141"
  selfLink: /apis/extensions/v1beta1/namespaces/pink/deployments/pink-carbon
  uid: 1d02f9e9-1d85-11e9-b012-06ed6a534df5
spec:
  progressDeadlineSeconds: 600
  replicas: 1
  revisionHistoryLimit: 10
  selector:
    matchLabels:
      app: pink-carbon

```



```
strategy:
  rollingUpdate:
    maxSurge: 25%
    maxUnavailable: 25%
  type: RollingUpdate
template:
  metadata:
    creationTimestamp: null
    labels:
      app: pink-carbon
  spec:
    containers:
      - env:
          - name: APP_NAMESPACE
            valueFrom:
              fieldRef:
                apiVersion: v1
                fieldPath: metadata.namespace
          - name: APP_NAME
            valueFrom:
              fieldRef:
                apiVersion: v1
                fieldPath: metadata.name
          - name: COLLECTOR_CONFIG
            valueFrom:
              configMapKeyRef:
                key: COLLECTOR_CONFIG
                name: pink-collector-config
          - name: INSTRUCTOR_CONFIG
            valueFrom:
              configMapKeyRef:
                key: INSTRUCTOR_CONFIG
                name: pink-collector-config
        image: ibmicpcoc/carbon:latest
        imagePullPolicy: Always
        name: pink-carbon
      resources:
        requests:
          cpu: "25"                                     <=== change value to 100m without quotes
          memory: 100Mi
        terminationMessagePath: /dev/termination-log
        terminationMessagePolicy: File
      dnsPolicy: ClusterFirst
      restartPolicy: Always
      schedulerName: default-scheduler
      securityContext: {}
      terminationGracePeriodSeconds: 30
    status:
      conditions:
        - lastTransitionTime: 2019-01-21T14:01:56Z
          lastUpdateTime: 2019-01-21T14:01:56Z
          message: Deployment does not have minimum availability.
          reason: MinimumReplicasUnavailable
          status: "False"
          type: Available
```

```
- lastTransitionTime: 2019-01-21T14:11:57Z
  lastUpdateTime: 2019-01-21T14:11:57Z
  message: ReplicaSet "pink-carbon-5c96bc649" has timed out progressing.
  reason: ProgressDeadlineExceeded
  status: "False"
  type: Progressing
observedGeneration: 1
replicas: 1
```

NOTE: You must save the file for the changes to take effect.

Result output:
deployment.extensions/pink-carbon edited

Did this resolve the issue?

```
Command to get pods in namespace:
oc -n <namespace> get pods
```

```
Example output:
NAME                                READY   STATUS    RESTARTS   AGE
pink-carbon-7784b95958-pct15       1/1     Running   0           2m
```

All references to "team" or <team> should be replaced with your team name which is the same as your namespace.

Use the debug flow to guide the steps you should attempt in diagnosis of the issue.

Resources

- K8 yaml - [doors.yaml](#)
- Dockerfile - [Dockerfile](#)

Useful information

Item	Value
cpu:	50m
memory:	50Mi
image:	ibmicpcoc/doors:latest
ports	none
Docker	CMD ["node", "app.js"]

Task description

Within your team namespace diagnose the pod that begins with <team>-doors

Task description
Use the label option -l app=<team>-doors when getting the pod status.
Download the resource K8 yaml file.
Use either of the delete-create-pod or edit-running-pod approaches to resolve the issue.
Did the pod deploy successfully? If not, correct the issue and re-create the K8 objects.

Check the "tag" of the image that is being pulled.

Diagnosis

Pod status

Command:

```
oc -n <your namespace> get pods -l app=<team>-doors      (replace <team> )
```

Example output:

NAME	READY	STATUS	RESTARTS	AGE
pink-doors-78b7f6598d-p8kvf	0/1	ImagePullBackOff	0	10m

Describe the pod (complete output from command is shown)

```
Name:          pink-doors-78b7f6598d-p8kvf
Namespace:     pink
Priority:       0
PriorityClassName: <none>
Node:          10.186.56.85/10.186.56.85
Start Time:    Mon, 21 Jan 2019 10:18:18 -0600
Labels:        app=pink-doors
               pod-template-hash=3349118043
. . .
      portions of output removed
. . .

Events:
  Type     Reason      Age           From          Message
  ----     -
  Normal    Scheduled    46s           default-scheduler    Successfully assigned pink/pink-doors-78b7f6598d-p8kvf to 10.186.56.85
  Normal    Pulling      28s (x2 over 43s)    kubelet, 10.186.56.85    pulling image "ibmicpcoc/doors:last"
  Warning   Failed       27s (x2 over 43s)    kubelet, 10.186.56.85    Failed to pull image "ibmicpcoc/doors:last": rpc error: code = Unknown desc = Error response from daemon: manifest for ibmicpcoc/doors:last not found
  Warning   Failed       27s (x2 over 43s)    kubelet, 10.186.56.85    Error: ErrImagePull
  Normal    BackOff      12s (x3 over 42s)    kubelet, 10.186.56.85    Back-off pulling image "ibmicpcoc/doors:last"
  Warning   Failed       12s (x3 over 42s)    kubelet, 10.186.56.85    Error: ImagePullBackOff
```

Multiple Warning messages are displayed in the Event section. Review all of the Warning messages.

In the "Events" section review the "Message" from the entry with "Type" Warning and "Reason" Failed

```
... Failed to pull image "ibmicpcoc/doors:last": rpc error: code = Unknown desc = Error response from daemon:
manifest for ibmicpcoc/doors:last not found
```

(output is from the first Failed message)

Problem discovered

The image cannot be located as indicated by the "Failed to pull image" message. The image tag last on the container is incorrect. The image tag should be latest.

Resolution

The edit-running-pod is shown in the following example to resolve the issue:

Command to edit the running pod:

```
KUBE_EDITOR="nano" oc -n <team> edit deployment/<team>-doors
```

Example is from the pink namespace. Modify the tag of the image to "latest"

```
# Please edit the object below. Lines beginning with a '#' will be ignored,
# and an empty file will abort the edit. If an error occurs while saving this file will be
# reopened with the relevant failures.
#
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  annotations:
    deployment.kubernetes.io/revision: "1"
  creationTimestamp: 2019-01-21T16:18:18Z
  generation: 1
  labels:
    app: pink-doors
  name: pink-doors
  namespace: pink
  resourceVersion: "5853628"
  selfLink: /apis/extensions/v1beta1/namespaces/pink/deployments/pink-doors
  uid: 29914949-1d98-11e9-b012-06ed6a534df5
spec:
  progressDeadlineSeconds: 600
  replicas: 1
  revisionHistoryLimit: 10
  selector:
    matchLabels:
      app: pink-doors
  strategy:
    rollingUpdate:
      maxSurge: 25%
      maxUnavailable: 25%
    type: RollingUpdate
```

```

template:
  metadata:
    creationTimestamp: null
    labels:
      app: pink-doors
  spec:
    containers:
      - env:
        - name: APP_NAMESPACE
          valueFrom:
            fieldRef:
              apiVersion: v1
              fieldPath: metadata.namespace
        - name: APP_NAME
          valueFrom:
            fieldRef:
              apiVersion: v1
              fieldPath: metadata.name
        - name: COLLECTOR_CONFIG
          valueFrom:
            configMapKeyRef:
              key: COLLECTOR_CONFIG
              name: pink-collector-config
        - name: INSTRUCTOR_CONFIG
          valueFrom:
            configMapKeyRef:
              key: INSTRUCTOR_CONFIG
              name: pink-collector-config
      image: ibmicpcoc/doors:last           <=== change the :last to :latest
      imagePullPolicy: Always

```

Ensure you have saved the modified file.

Result output:
deployment/pink-doors

Validate the pod status is Running.

```

Command:
  oc -n <team> get pods

Example output:

```

NAME	READY	STATUS	RESTARTS	AGE
pink-doors-767f49c748-6gvvg	1/1	Running	0	1m

All references to "team" or <team> should be replaced with your team name which is the same as your namespace.

Use the debug flow to guide the steps you should attempt in diagnosis of the issue.

Resources

- K8 yaml - [avail.yaml](#)
- Dockerfile - [Dockerfile](#)

Useful information

Item	Value
cpu:	50m
memory:	50Mi
image:	ibmicpcoc/avail:latest
ports	none
Run K8 spec	command: ["/bin/bash", "-c", "/app/avail.sh"]

Task description
Within the "avail" namespace research the pod that begins with "avail".
Why is the pod not deploying?
Review K8 definitions for controlling privileges e.g. PSP, RoleBinding, Roles etc.
Download the resource K8 yaml file.
Edit the file replacing <team> with your team name.
Create the K8 objects.

What rolebinding is defined for avail namespace?
What rolebinding is defined for <team> namespace?
Review the clusterroles for the cluster.
Reivew the pod security policies for the cluster.

Diagnosis

Command to check pods in namespace:

```
oc -n avail get pods
```

Example output:

NAME	READY	STATUS	RESTARTS	AGE
avail-all-65b8448469-rqt5g	0/1	CreateContainerConfigError	0	1d

Command to describe the selected pod in the namespace:

```
oc -n avail describe pod avail-all-65b8448469-rqt5g
```

Example output:

```
Name:          avail-all-65b8448469-rqt5g
Namespace:     avail
Priority:       0
PriorityClassName: <none>
Node:          10.186.56.85/10.186.56.85
Start Time:    Sat, 19 Jan 2019 13:57:24 -0600
Labels:        app=avail-all
                pod-template-hash=2164004025
```

. . . < portions of the describe output not shown> . . .

Events:

Type	Reason	Age	From	Message
Normal	Scheduled	28m	default-scheduler	Successfully assigned avail/avail-698964bc87-5k8vf to 10.186.56.85
Normal	Pulled	26m (x8 over 28m)	kubelet, 10.186.56.85	Successfully pulled image "avail"
Warning	Failed	26m (x8 over 28m)	kubelet, 10.186.56.85	Error: container has runAsNonRoot and image will run as root

In the "Events" section review the "Message" from the entry with "Type" Warning and "Reason" Failed

```
... Error: container has runAsNonRoot and image will run as root
```

What rolebinding are defined for the **avail** namespace?

```
Command to check rolebindings:
oc get rolebinding -n avail
```

Example output:

```
No resources found.
```

Compare rolebindings for your **team** namespace.

```
Command to check rolebindings:
oc get rolebinding -n <team>
```

Example output:

NAME	AGE	ROLE	USERS	GROUPS
SERVICEACCOUNTS				
ibm-privileged-clusterrole-rolebinding	16h	ClusterRole/ibm-privileged-clusterrole		
system:serviceaccounts:aqua				

Review the clusterrole definitions for the cluster.

```
Command to view clusterrole
k get clusterrole
```

Example output:

```
NAME                                     AGE
admin                                   17h
cluster-admin                           17h
edit                                     17h
extension                               17h
ibm-anyuid-clusterrole                   17h
ibm-anyuid-hostaccess-clusterrole        17h
ibm-anyuid-hostpath-clusterrole          17h
ibm-cert-manager-cert-manager            17h
ibm-privileged-clusterrole               17h
ibm-restricted-clusterrole               17h
icp-admin-aggregate                      17h
icp-edit-aggregate                       17h
icp-operate-aggregate                    17h
icp-view-aggregate                       17h

. . . data truncated . . .
```

Describe the clusterrole for ibm-privileged-clusterrole

Command to describe:

```
oc describe clusterrole ibm-privileged-clusterrole
```

Example output;

```
Name:      ibm-privileged-clusterrole
Labels:    <none>
Annotations:  oc.kubernetes.io/last-applied-configuration=
{"apiVersion":"rbac.authorization.k8s.io/v1","kind":"ClusterRole","metadata":{"annotations":{},"name":"ibm-privileged-clusterrole","namespace":""},"rules":...
PolicyRule:
  Resources          Non-Resource URLs  Resource Names      Verbs
  -----
  podsecuritypolicies.extensions  []                  [ibm-privileged-psp]  [use]
```

Review the Pod Security Policies.

Command to view Pod Security Policy:

```
oc get psp
```

Example output:

```
NAME          RUNASUSER  PRIV  CAPS  SUPGROUP  READONLYROOTFS  VOLUMES
```



```

ibm-anyuid-hostaccess-psp    false
SETPCAP,AUDIT_WRITE,CHOWN,NET_RAW,DAC_OVERRIDE,FOWNER,FSETID,KILL,SETUID,SETGID,NET_BIND_SERVICE,SYS_CHROOT,SETFCAP
RunAsAny    RunAsAny            RunAsAny    RunAsAny    false        *
ibm-anyuid-hostpath-psp     false
SETPCAP,AUDIT_WRITE,CHOWN,NET_RAW,DAC_OVERRIDE,FOWNER,FSETID,KILL,SETUID,SETGID,NET_BIND_SERVICE,SYS_CHROOT,SETFCAP
RunAsAny    RunAsAny            RunAsAny    RunAsAny    false        *
ibm-anyuid-psp              false
SETPCAP,AUDIT_WRITE,CHOWN,NET_RAW,DAC_OVERRIDE,FOWNER,FSETID,KILL,SETUID,SETGID,NET_BIND_SERVICE,SYS_CHROOT,SETFCAP
RunAsAny    RunAsAny            RunAsAny    RunAsAny    false
configMap,emptyDir,projected,secret,downwardAPI,persistentVolumeClaim
ibm-privileged-psp         true        *
RunAsAny    RunAsAny            RunAsAny    RunAsAny    false        *
ibm-restricted-psp         false
RunAsAny    MustRunAsNonRoot    MustRunAs    MustRunAs    false
configMap,emptyDir,projected,secret,downwardAPI,persistentVolumeClaim

```

Problem discovered

The "avail" namespace does not have the proper authority to run the "avail" pod. The avail pod must be deployed within a namespace that has the proper authority. Your team namespace has the proper authority.

Resolution

Download the K8 Yaml file from the resources section and save locally. Once saved, edit the file and change the namespace metadata parameter in the file and deploy the pod.

```

Example saved file avail.yaml (only a portion of file is shown below)

apiVersion: apps/v1
kind: Deployment
metadata:
  name: avail
  namespace: <team>      # change <team> to your namespace and save the file

----
Command to create the new pod:
  oc create -f avail.yaml

Result output:
  deployment.apps/avail created

----
Verify issue is resolved. Pod status should be "Running":

Command to get pods in namespace:
  oc -n <team> get pods      # change <team> to your team namespace

Example output:
  avail-698964bc87-2fpw8    1/1      Running    0      1m

```

All references to "team" or <team> should be replaced with your team name which is the same as your namespace.

Resources

- K8 yaml - [eagle.yaml](#)
- Dockerfile - [Dockerfile](#)

Useful information

Item	Value
cpu:	50m
memory:	50Mi
image:	ibmicpcoc/eagle:latest
ports	4100
Docker	CMD ["node", "server.js"]

Task description

This lab uses the pod with a name that starts with **<team>-eagle**

The web application is not working properly. The application is has a K8 Deployment and Service defined.

Research why the web application is not working properly.

Once you have resolved the issue locate the NodePort (is a number in the 30000 range) for the service. Example: `oc get svc -n <team> -o wide`

Using the same IP that has been used to access the Collector now access the the web application using the newly located node port number. Example url to access web application: <http://xxx.xxx.xxx.xxx:NodePort>

Once the web application is successfully accessed press the button to complete the lab.

-
- All exposed port definitions must match.
 - What port should the application be available on? Refer to useful information.

Diagnosis

The pod is running successfully yet describing the pod can provide information about the configured K8 objects. Describe the pod that begins with: <team>-eagle

```
Commad to get pods in namespace
oc -n <team> get pods           # Replace <team> with namespace name

Command to describe the pod           # Use the pod name from the previous output
oc -n <team> describe pod <pod>      # Replace <team> with namespace name
```

Review the port definitions from the describe output

Show something here

Problem discovered

The ports do not match for the Deployment and Service definitions.

Resolution

Edit the Service definition and change the port from 4010 to 4100.

Add detailed steps here

All references to "team" or <team> should be replaced with your team name which is the same as your namespace.

Resources

- K8 yaml - [floor.yaml](#)
- Dockerfile - [Dockerfile](#)

Useful information

Item	Value
cpu:	50m
memory:	50Mi
image:	ibmicpcoc/floor:latest
ports	none
YAML	command: ["node", "app.js"]

Task description

A container wihtin a successfully deployed pod is not working properly. Research the running container to diagnose the issue.

View the logs of the running container.

Correct the issue inside the running container.

- Exec into the running container
- Use touch, nano, or echo with piping to assit in resolving the issue

Diagnosis

Check the logs of the running container that begins with <team>

Command to get pods in namespace
oc -n <team> get pods <=== Replace <team>

Example output from "pink" namespace

NAME	READY	STATUS	RESTARTS	AGE
pink-floor-6ff9f54f44-zpchp	1/1	Running	0	41s

Get the logs for the pod

```
oc -n <team> logs -f <pod>          <=== Replace <team> and <pod>
                                     Use the pod name from the get pods result
```

Instructions from viewing the log

```
1/21/2019, 10:21:14 PM :: clnt012i - Check for file: /app/team.txt check count: 43
1/21/2019, 10:21:14 PM :: clnt013i - The file team.txt in the /app directory must exist for this lab to be
completed.
1/21/2019, 10:21:14 PM :: clnt014i - Create the file in the running container.
```

Problem discovered

The file team.txt is missing from the /app directory in the running container.

Resolution

Two methods can be used to resolve of creating the file.

First method is to run a "command" using the oc CLI from outside the container.

Command to get pods in namespace

```
oc -n <team> get pods          <=== Replace <team>
```

Example output from "pink" namespace

NAME	READY	STATUS	RESTARTS	AGE
pink-floor-6ff9f54f44-zpchp	1/1	Running	0	41s

Add the team.txt file using the touch command from outside the container.

```
oc exec -n pink pink-floor-6ff9f54f44-zpchp -- sh -c "touch /app/team.txt"
```

The above command is using 'sh'. The 'sh' capability must be installed in the container for this to work.

Example result output: (wait a few seconds for the messages to show)

```
1/21/2019, 10:25:30 PM :: clnt014i - Create the file in the running container.
1/21/2019, 10:25:45 PM :: -----
1/21/2019, 10:25:45 PM :: clnt008i - File located. Reporting to collector.
1/21/2019, 10:25:45 PM :: -----
1/21/2019, 10:25:45 PM :: clnt007i - Student count: 61 from /pink/pink-floor-6ff9f54f44-zpchp
1/21/2019, 10:25:45 PM :: clnt010i - Instructor count: 1 from /pink/pink-floor-6ff9f54f44-
```

The clnt007i and clnt010i messages are produced once the file has been loacted.

Second method is to exec into the running container and create the file from a shell prompt. This method requires 'sh' capability must be installed in the container for this to work.

Command to get pods in namespace

```
oc -n <team> get pods
```

<=== Replace <team>

Example output from "pink" namespace

NAME	READY	STATUS	RESTARTS	AGE
pink-floor-6ff9f54f44-zpchp	1/1	Running	0	41s

Open a terminal session with the running session

Add the team.txt file using the touch command from outside the container.

```
oc exec -it -n pink pink-floor-6ff9f54f44-zpchp -- sh
```

The above command is using 'sh'. The 'sh' capability must be installed in the container for this to work.

Example result output:

```
/app #
```

Create the file using touch by entering the following command:

```
touch team.txt
```

Notice the "/app" directory is not included as part of the touch command since the prompt is open to that directory.

Example result output: (wait a few seconds for the messages to show)

```
1/21/2019, 10:25:30 PM :: clnt014i - Create the file in the running container.
```

```
1/21/2019, 10:25:45 PM :: -----
```

```
1/21/2019, 10:25:45 PM :: clnt008i - File located. Reporting to collector.
```

```
1/21/2019, 10:25:45 PM :: -----
```

```
1/21/2019, 10:25:45 PM :: clnt007i - Student count: 61 from /pink/pink-floor-6ff9f54f44-zpchp
```

```
1/21/2019, 10:25:45 PM :: clnt010i - Instructor count: 1 from /pink/pink-floor-6ff9f54f44-
```

The clnt007i and clnt010i messages are produced once the file has been loacted.

All references to "team" or <team> should be replaced with your team name which is the same as your namespace.

Resources

- K8 yaml - [gonzo.yaml](#)
- Dockerfile - [Dockerfile](#)

Useful information

Item	Value
cpu:	50m
memory:	50Mi
image:	ibmicpcoc/gonzo:latest
ports	none
YAML	command: ["/bin/bash", "-c", "/app/gonzo.sh"]

Task description
A pod that begins with <team>-gonzo is failing creation.
Research the issue to determine what is causing the failure.
Edit the gonzo.yaml file to correct the issue.
Verify the deployment successfully deployed

- What ENTRYPOINT or CMD is defined for the Docker image?
- What container "command" parameter is defined for the pod definition?
- Command: `docker history ibmicpcoc/gonzo --no-trunc` can also be used to check the docker image.
- The gonzo.yaml must be modified to correct the issue. You will not be allowed to rebuild or modify the Docker image.

Diagnosis

Command to get pods in namespace

```
oc -n <team> get pods
```

<=== Replace <team>

Example output from "pink" namespace

NAME	READY	STATUS	RESTARTS	AGE
pink-gonzo-75d79787b7-88pnr	0/1	CrashLoopBackOff	4	2m

Command to describe pod that is failing. Following example using above pod and pink namespace.

```
oc describe pod pink-gonzo-75d79787b7-88pnr -n pink
```

Example output:

```
Name:                pink-gonzo-75d79787b7-88pnr
Namespace:           pink
Priority:              0
PriorityClassName:    <none>
Node:                10.186.56.85/10.186.56.85
Start Time:          Mon, 21 Jan 2019 18:13:15 -0600
Labels:               app=pink-gonzo
                     pod-template-hash=3183534363
```

. . .

portions of output removed

. . .

```

Conditions:
  Type           Status
  Initialized     True
  Ready           False
  ContainersReady False
  PodScheduled    True

Volumes:
  default-token-mq64m:
    Type:          Secret (a volume populated by a Secret)
    SecretName:    default-token-mq64m
    Optional:      false
QoS Class:       Burstable
Node-Selectors:  <none>
Tolerations:     node.kubernetes.io/memory-pressure:NoSchedule
                  node.kubernetes.io/not-ready:NoExecute for 300s
                  node.kubernetes.io/unreachable:NoExecute for 300s

Events:
  Type     Reason      Age           From          Message
  ----     -
  Normal   Scheduled   11m          default-scheduler   Successfully assigned pink/pink-gonzo-75d79787b7-88pnr to 10.186.56.85
  Normal   Created     10m (x4 over 11m) kubelet, 10.186.56.85   Created container
  Normal   Started     10m (x4 over 11m) kubelet, 10.186.56.85   Started container
  Normal   Pulling     9m (x5 over 11m) kubelet, 10.186.56.85   pulling image "ibmicpcoc/gonzo:latest"
  Normal   Pulled      9m (x5 over 11m) kubelet, 10.186.56.85   Successfully pulled image "ibmicpcoc/gonzo:latest"
  Warning  BackOff     58s (x46 over 11m) kubelet, 10.186.56.85   Back-off restarting failed container

```

In the "Events" section review the "Message" from the entry with "Type" Warning and "Reason" BackOff

```
... Back-off restarting failed container
```

Check the image for the command or entrypoint defined to execute when the container is created

Review the Dockerfile provided in the Resources section of this lab.

Browse the Dockerfile

Click the Dockerfile link in resource section and review the entrypoint or command defined to start when container is created.

(or)

Check the Docker image

```
docker history ibmicpcoc/gonzo --no-trunc
```

Problem discovered

The container is ending as soon as it starts. The entrypoint or command that executes when the container starts is not defined in either the Dockerfile or gonzo.yaml file.

Resolution

Add the "command" parameter to the pod container definition using the file gonzo.yaml provided in the Resources section of this lab. The "command" parameter should start the bash script /app/gonzo.sh using /bin/bash

```
command: ["/bin/bash", "-c", "/app/gonzo.sh"]
```

```
Add the "command" parameter to the container:apiVersion: apps/v1kind: Deploymentmetadata: name: pink-gonzo
namespace: pink labels: app: pink-gonzospec: selector: matchLabels: app: pink-gonzo replicas: 1
template: metadata: labels: app: pink-gonzo spec: containers: - name: pink-gonzo
image: ibmicpcoc/gonzo:latest imagePullPolicy: Always command: ["/bin/bash", "-c", "/app/gonzo.sh"]
<=== insert this line. . . remainder of file not shown . . .Save the modified fileCommand to delete the current
deployed pod oc -n <team> delete -f gonzo.yamlExample output: deployment.apps/pink-gonzo deleteCommand to
deploy the updated pod oc -n <team> create -f gonzo.yamlExample output: deployment.apps/pink-gonzo
createdCommand to verify the updated pod is running oc -n <team> get podsExample output: NAME
READY STATUS RESTARTS AGE pink-gonzo-67834787b7-234xy 1/1 Running 0 2m
```

All references to "team" or <team> should be replaced with your team name which is the same as your namespace.

Resources

- K8 yaml - [igloo.yaml](#)
- Dockerfile - [Dockerfile](#)

Useful information

Item	Value
cpu:	50m
memory:	50Mi
image:	ibmicpcoc/igloo:latest
ports	none
YAML	command: ["node", "app.js"]
Misc	Application waits

Task description
A pod that begins with <team>-igloo is frequently restarting.
Research the issue to determine what is causing the pod to restart frequently.
Reiview the pod log to determine how long the http server waits to be started.
Edit the igloo.yaml file to correct the issue.
Verify the deployment successfully deployed.
Get the NodePort for the red-igloo service.
Get the IP address for the master node.
Using the above NodePort and the master ip address access the url: http://:

- How long do both probes delay before starting?

Diagnosis

Command to get pods in namespace

```
oc -n <team> get pods
```

<=== Replace <team>

Example output from "red" namespace

NAME	READY	STATUS	RESTARTS	AGE
red-igloo-7b85976d87-x6z6r	0/1	Running	3	2m

Command to view the pod details

```
oc describe po red-igloo-7b85976d87-x6z6r
```

Name: red-igloo-7b85976d87-x6z6r

Namespace: red

Priority: 0

PriorityClassName: <none>

Node: gfstst.169.62.225.201.nip.io/169.62.225.201

Start Time: Tue, 03 Sep 2019 20:06:13 -0400

Labels: app=red-igloo

pod-template-hash=3641532843

Annotations: openshift.io/scc=restricted

Status: Running

IP: 10.129.0.94

Controlled By: ReplicaSet/red-igloo-7b85976d87

Containers:

red-igloo:

Container ID: docker://e9b6049395fa281c1ca0d6e63001ac3226fc211c5948bf1673023c9dc6f74f37

Image: ibmicpcoc/igloo:latest

Image ID: docker-

pullable://docker.io/ibmicpcoc/igloo@sha256:4968f5c1ca641e3267d9a163c68eceb307973e06a30df51a47d86dcd0e301a40

Port: <none>

Host Port: <none>

State: Running

Started: Tue, 03 Sep 2019 20:06:49 -0400

Last State: Terminated

Reason: Error

Exit Code: 137

Started: Tue, 03 Sep 2019 20:06:16 -0400

Finished: Tue, 03 Sep 2019 20:06:48 -0400

Ready: False

Restart Count: 1

Requests:

cpu: 50m

memory: 50Mi

Liveness: http-get http://:4100/health delay=1s timeout=1s period=2s #success=1 #failure=1

Readiness: http-get http://:4100/ready delay=1s timeout=1s period=5s #success=1 #failure=3

Environment:

APP_NAMESPACE: red (v1:metadata.namespace)

APP_NAME: red-igloo-7b85976d87-x6z6r (v1:metadata.name)

COLLECTOR_CONFIG: <set to the key 'COLLECTOR_CONFIG' of config map 'red-collector-config'> Optional:

false

```

    INSTRUCTOR_CONFIG: <set to the key 'INSTRUCTOR_CONFIG' of config map 'red-collector-config'> Optional:
false
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-dxnzt (ro)
    Conditions:
      Type              Status
      Initialized        True
      Ready              False
      ContainersReady    False
      PodScheduled       True
    Volumes:
      default-token-dxnzt:
        Type:          Secret (a volume populated by a Secret)
        SecretName:    default-token-dxnzt
        Optional:      false
    QoS Class:          Burstable
    Node-Selectors:     node-role.kubernetes.io/compute=true
    Tolerations:        node.kubernetes.io/memory-pressure:NoSchedule
    Events:
      Type    Reason    Age    From    Message
      ----    -
      Normal  Scheduled  42s    default-scheduler    Successfully assigned red/red-igloo-7b85976d87-x6z6r to gfstst.169.62.225.201.nip.io
      Normal  Pulling    7s (x2 over 40s)    kubelet, gfstst.169.62.225.201.nip.io    pulling image "ibmicpcoc/igloo:latest"
      Normal  Killing    7s    kubelet, gfstst.169.62.225.201.nip.io    Killing container with id docker://red-igloo:Container failed liveness probe.. Container will be killed and recreated.
      Normal  Pulled     6s (x2 over 39s)    kubelet, gfstst.169.62.225.201.nip.io    Successfully pulled image "ibmicpcoc/igloo:latest"
      Normal  Created    6s (x2 over 39s)    kubelet, gfstst.169.62.225.201.nip.io    Created container
      Normal  Started    6s (x2 over 39s)    kubelet, gfstst.169.62.225.201.nip.io    Started container
      Warning  Unhealthy  4s (x2 over 38s)    kubelet, gfstst.169.62.225.201.nip.io    Liveness probe failed: Get http://10.129.0.94:4100/health: dial tcp 10.129.0.94:4100: connect: connection refused
      Warning  Unhealthy  2s (x3 over 37s)    kubelet, gfstst.169.62.225.201.nip.io    Readiness probe failed: Get http://10.129.0.94:4100/ready: dial tcp 10.129.0.94:4100: connect: connection refused

```

Command to view the pod logs

```
oc logs red-igloo-5dd5b6c7b8-jqdvr
```

Example output

```
9/4/2019, 1:54:54 AM :: igloo900i - Waiting 10 seconds to start HTTP server
```

All references to "team" or <team> should be replaced with your team name which is the same as your namespace.

Resources

- K8 yaml - [jazzy.yaml](#)
- Dockerfile - [Dockerfile](#)

Useful information

Item	Value
------	-------

Item	Value
cpu:	50m
memory:	50Mi
image:	ibmicpcoc/jazzy:latest
ports	9000
YAML	command: ["node", "app.js"]
Misc	Application waits

Task description
A pod that begins with <team>-jazzy is frequently restarting.
Research the issue to determine what is causing the pod to restart frequently.
Reiview the pod log to determine how long the application http server waits to be started.
Edit the igloo.yaml file to correct the issue.
Verify the deployment successfully deployed.
Get the NodePort for the red-igloo service.
Get the IP address for the master node.
Using the above NodePort and the master ip address access the url: http://:

- How long do Readiness and Liveness probes delay before starting?

Diagnosis

Command to get pods in namespace

```
oc -n <team> get pods
```

<=== Replace <team>

Example output from "red" namespace

NAME	READY	STATUS	RESTARTS	AGE
red-igloo-7b85976d87-x6z6r	0/1	Running	3	2m

Command to view the pod details

```
oc describe po red-igloo-7b85976d87-x6z6r
```

```
Name: red-igloo-7b85976d87-x6z6r
Namespace: red
Priority: 0
PriorityClassName: <none>
Node: gfstst.169.62.225.201.nip.io/169.62.225.201
Start Time: Tue, 03 Sep 2019 20:06:13 -0400
Labels: app=red-igloo
        pod-template-hash=3641532843
Annotations: openshift.io/scc=restricted
Status: Running
```

```

IP: 10.129.0.94
Controlled By: ReplicaSet/red-igloo-7b85976d87
Containers:
  red-igloo:
    Container ID: docker://e9b6049395fa281c1ca0d6e63001ac3226fc211c5948bf1673023c9dc6f74f37
    Image: ibmicpcoc/igloo:latest
    Image ID: docker-
pullable://docker.io/ibmicpcoc/igloo@sha256:4968f5c1ca641e3267d9a163c68eceb307973e06a30df51a47d86dcd0e301a40
    Port: <none>
    Host Port: <none>
    State: Running
      Started: Tue, 03 Sep 2019 20:06:49 -0400
    Last State: Terminated
      Reason: Error
      Exit Code: 137
      Started: Tue, 03 Sep 2019 20:06:16 -0400
      Finished: Tue, 03 Sep 2019 20:06:48 -0400
    Ready: False
    Restart Count: 1
    Requests:
      cpu: 50m
      memory: 50Mi
    Liveness: http-get http://:4100/health delay=1s timeout=1s period=2s #success=1 #failure=1
    Readiness: http-get http://:4100/ready delay=1s timeout=1s period=5s #success=1 #failure=3
    Environment:
      APP_NAMESPACE: red (v1:metadata.namespace)
      APP_NAME: red-igloo-7b85976d87-x6z6r (v1:metadata.name)
      COLLECTOR_CONFIG: <set to the key 'COLLECTOR_CONFIG' of config map 'red-collector-config'> Optional:
false
      INSTRUCTOR_CONFIG: <set to the key 'INSTRUCTOR_CONFIG' of config map 'red-collector-config'> Optional:
false
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-dxnzt (ro)
    Conditions:
      Type          Status
      Initialized    True
      Ready          False
      ContainersReady False
      PodScheduled   True
    Volumes:
      default-token-dxnzt:
        Type: Secret (a volume populated by a Secret)
        SecretName: default-token-dxnzt
        Optional: false
    QoS Class: Burstable
    Node-Selectors: node-role.kubernetes.io/compute=true
    Tolerations: node.kubernetes.io/memory-pressure:NoSchedule
    Events:
      Type      Reason      Age      From      Message
      ----      -
      Normal    Scheduled   42s      default-scheduler      Successfully assigned red/red-igloo-7b85976d87-x6z6r to gfstst.169.62.225.201.nip.io
      Normal    Pulling     7s (x2 over 40s)      kubelet, gfstst.169.62.225.201.nip.io pulling image "ibmicpcoc/igloo:latest"
      Normal    Killing     7s      kubelet, gfstst.169.62.225.201.nip.io Killing container with id

```

```

docker://red-igloo:Container failed liveness probe.. Container will be killed and recreated.
Normal    Pulled      6s (x2 over 39s)  kubelet, gfstst.169.62.225.201.nip.io  Successfully pulled image
"ibmicpcoc/igloo:latest"
Normal    Created     6s (x2 over 39s)  kubelet, gfstst.169.62.225.201.nip.io  Created container
Normal    Started     6s (x2 over 39s)  kubelet, gfstst.169.62.225.201.nip.io  Started container
Warning   Unhealthy   4s (x2 over 38s)  kubelet, gfstst.169.62.225.201.nip.io  Liveness probe failed: Get
http://10.129.0.94:4100/health: dial tcp 10.129.0.94:4100: connect: connection refused
Warning   Unhealthy   2s (x3 over 37s)  kubelet, gfstst.169.62.225.201.nip.io  Readiness probe failed: Get
http://10.129.0.94:4100/ready: dial tcp 10.129.0.94:4100: connect: connection refused

```

Checking the running pod for application information regarding the startup delay.

```

Command to view the pod logs
oc logs red-igloo-5dd5b6c7b8-jqdvr

Example output
9/4/2019, 1:54:54 AM :: igloo900i - Waiting 10 seconds to start HTTP server

```

Problem discovered

The Readiness and Liveness probes do not delay long enough to allow the application to start.

Resolution

Modify the ReadinessProbe initialDelaySeconds to be longer than the ten seconds the application takes to start. Also modify the LivenessProbe initialDelaySeconds and periodSeconds to be longer than the ten seconds the application takes to start.

All references to "team" or <team> should be replaced with your team name which is the same as your namespace.

Resources

- K8 yaml - [lacey.yaml](#)
- Dockerfile - [Dockerfile](#)

Useful information

Item	Value
cpu:	50m
memory:	50Mi
image:	ibmicpcoc/lacey:latest
ports	none
YAML	command: ["node", "app.js"]
Misc	Application waits

Task description

A pod that begins with <team>-igloo is frequently restarting.

Research the issue to determine what is causing the pod to restart frequently.

Task description
Reiview the pod log to determine how long the application http server waits to be started.
Edit the igloo.yaml file to correct the issue.
Verify the deployment successfully deployed.
Get the NodePort for the red-igloo service.
Get the IP address for the master node.
Using the above NodePort and the master ip address access the url: http://:

Provide a hint.

Diagnosis

Checking the running pod for application information.

```
Command to view the pod logs
oc logs red-igloo-5dd5b6c7b8-jqdv

Example output
9/4/2019, 1:54:54 AM :: igloo900i - Waiting 10 seconds to start HTTP server
```

Problem discovered

Describe the problem.

Resolution

Describe the resolution.

All references to "team" or <team> should be replaced with your team name which is the same as your namespace.

Resources

- K8 yaml - [magma.yaml](#)
- Dockerfile - [Dockerfile](#)

Useful information

Item	Value
cpu:	50m
memory:	50Mi
image:	ibmicpcoc/magma:latest
ports	none
YAML	command: ["node", "app.js"]

Secret Parameter	Value
------------------	-------

Secret Parameter	Value
Name	<team>-secret-file
Content	Base64 encoded: debug me
Type	Opaque
Mount	/var/config
File	secret.txt

ConfigMap Parameter	Value
Name	<team>-configmap-file
Content	debug
Mount	/var/secret
File	config.txt

Task description
A pod that begins with <team>-magma has a status of ContainerCreating.
Research the issue to determine what is causing the pod to be in this status.
Describe the pod to assist in determining why this issue is occurring.
Reiview the pod log to determine how long the application http server waits to be started.
Edit the magma.yaml file to correct the issue.
Verify the deployment successfully deployed.

Create the secret and configmap.

Diagnosis

Checking the running pod for information.

Command to view pod status

```
oc get pods -n <team>
```

Example output

NAME	READY	STATUS	RESTARTS	AGE
red-magma-6c4b56dbc9-kdtkv	0/1	ContainerCreating	0	11s

Command to describe the pod

```
oc describe po red-magma-6c4b56dbc9-kdtkv
```

Example output

```
Name: red-magma-6c4b56dbc9-kdtkv
```

```

Namespace:      red
Priority:        0
PriorityClassName: <none>
Node:           gfstst.169.62.225.207.nip.io/169.62.225.207
Start Time:     Sat, 07 Sep 2019 12:48:53 -0400
Labels:         app=red-magma
                pod-template-hash=2706128675
Annotations:    openshift.io/scc=restricted
Status:         Pending
IP:
Controlled By:  ReplicaSet/red-magma-6c4b56dbc9
Containers:
  red-magma:
    Container ID:
    Image:         ibmicpcoc/magma:latest
    Image ID:
    Port:          <none>
    Host Port:     <none>
    Command:
      node
      app.js
    State:         Waiting
      Reason:      ContainerCreating
    Ready:         False
    Restart Count: 0
    Requests:
      cpu:          50m
      memory:       50Mi
    Environment:
      APP_NAMESPACE:      red (v1:metadata.namespace)
      APP_NAME:            red-magma-6c4b56dbc9-kdtkv (v1:metadata.name)
      COLLECTOR_CONFIG:   <set to the key 'COLLECTOR_CONFIG' of config map 'red-collector-config'> Optional:
false
      INSTRUCTOR_CONFIG:  <set to the key 'INSTRUCTOR_CONFIG' of config map 'red-collector-config'> Optional:
false
    Mounts:
      /var/config from configvol (rw)
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-dxnzt (ro)
      /var/secret from secretvol (rw)
    Conditions:
      Type             Status
    Initialized        True
    Ready              False
    ContainersReady    False
    PodScheduled       True
  Volumes:
    configvol:
      Type:          ConfigMap (a volume populated by a ConfigMap)
      Name:          red-configmap-file
      Optional:      false
    secretvol:
      Type:          Secret (a volume populated by a Secret)
      SecretName:    red-secret-file
      Optional:      false
    default-token-dxnzt:

```



```

Type:          Secret (a volume populated by a Secret)
SecretName:    default-token-dxnzt
Optional:      false
QoS Class:     Burstable
Node-Selectors: node-role.kubernetes.io/compute=true
Tolerations:   node.kubernetes.io/memory-pressure:NoSchedule
Events:
  Type    Reason      Age    From          Message
  ----    -
  Normal  Scheduled   42s    default-scheduler  Successfully assigned red/red-
magma-6c4b56dbc9-kdtkv to gfstst.169.62.225.207.nip.io
  Warning FailedMount 10s (x7 over 42s) kubelet, gfstst.169.62.225.207.nip.io MountVolume.SetUp failed for
volume "configvol" : configmaps "red-configmap-file" not found
  Warning FailedMount 10s (x7 over 42s) kubelet, gfstst.169.62.225.207.nip.io MountVolume.SetUp failed for
volume "secretvol" : secrets "red-secret-file" not found

```

Problem discovered

Two volume mounts are failing for configvol and secretvol volumes. These mounts require a configmap and secret definitions that are not found.

Resolution

Create a secret of opaque type with base64 encoded value 'debug me' without the quotes. This secret is accessed via a volume

oc create secret generic apikey --from-file=./apikey.txtsecret "apikey" created

All references to "team" or <team> should be replaced with your team name which is the same as your namespace.

Desired environment:

Deployment of an application that uses persistent storage. The storage is implemented as static storage with a PV and PVC. The PV uses NFS based storage.

Note: **This lab requires the student to resolve multiple issues**

Resources

- K8 yaml - [offer.yaml](#)
- Dockerfile - [Dockerfile](#)

Useful information

Item	Value
cpu:	50m
memory:	50Mi
image:	ibmicpcoc/offer:latest
ports	none
YAML	command: ["node", "app.js"]

PV Parm	Value
---------	-------

PV Parm	Value
metadata.name	<team>-pv
metadata.labels.user	<team>
spec.capacity.storage	1Mi
spec.accessModes	ReadWriteOnce
spec.nfs.path	/storage/<team>/pvc001
spec.nfs.server	
persistentVolumeReclaimPolicy	recycle

PVC Parm	Value
metadata.name	<team>-offer
metadata.namespace	<team>
spec.resources.requests.storage	1Mi
spec.accessModes	ReadWriteOnce
spec.selector.matchLabels.user	<team>

Task description
Research multiple issues that are preventing the succesful deployment of the pod.
During the debugging be sure to describe the pod and view the pod logs.
Edit the offer.yaml file to correct the issues. (repeat)
Verify the deployment successfully deployed.

You have admin rights to create directories.

Diagnosis 1

Checking the running pod for application information.

```
Command to describe the pod
oc describe po red-offer-6cdf4749df-rtfwg
```

Example output

```
Name:                red-offer-6cdf4749df-rtfwg
Namespace:           red
Priority:             0
PriorityClassName:    <none>
Node:                <none>
Labels:              app=red-offer
```

```

pod-template-hash=2789030589
Annotations:      openshift.io/scc=restricted
Status:           Pending
IP:
Controlled By:    ReplicaSet/red-offer-6cdf4749df
Containers:
  red-offer:
    Image:          ibmicpcoc/offer:latest
    Port:           <none>
    Host Port:      <none>
    Command:
      node
      app.js
    Requests:
      cpu:          50m
      memory:       50Mi
    Environment:
      APP_NAMESPACE:  red (v1:metadata.namespace)
      APP_NAME:       red-offer-6cdf4749df-rtfwg (v1:metadata.name)
      COLLECTOR_CONFIG: <set to the key 'COLLECTOR_CONFIG' of config map 'red-collector-config'> Optional:
false
      INSTRUCTOR_CONFIG: <set to the key 'INSTRUCTOR_CONFIG' of config map 'red-collector-config'> Optional:
false
    Mounts:
      /data from offer-data (rw)
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-dxnzt (ro)
    Conditions:
      Type          Status
      PodScheduled  False
    Volumes:
      offer-data:
        Type:          PersistentVolumeClaim (a reference to a PersistentVolumeClaim in the same namespace)
        ClaimName:     red-offr
        ReadOnly:      false
      default-token-dxnzt:
        Type:          Secret (a volume populated by a Secret)
        SecretName:    default-token-dxnzt
        Optional:      false
    QoS Class:       Burstable
    Node-Selectors:  node-role.kubernetes.io/compute=true
    Tolerations:     node.kubernetes.io/memory-pressure:NoSchedule
    Events:
      Type          Reason          Age          From          Message
      ----          -
      Warning       FailedScheduling  1m (x25 over 2m)  default-scheduler  persistentvolumeclaim "red-offr" not found

```

Problem 1 discovered

Message from the describe indicates the PVC is not found. The Deployment spec.template.spec.volumes.persistentVolumeClaim.claimName does not match the name of the defined PVC.

Resolution 1

Change the `spec.template.spec.volumes.persistentVolumeClaim.claimName` parameter to match the name of the defined PVC: `<team>-offer` (notice the dash between `<team>` and `offer`)

Diagnosis 2

Checking the running pod for application information.

Command to view the pod logs

```
oc describe po red-offer-66c6488b56-m9wfr
```

Example output

```
Name:                red-offer-66c6488b56-m9wfr
Namespace:           red
Priority:             0
PriorityClassName:    <none>
Node:                gfstst.169.62.225.201.nip.io/169.62.225.201
Start Time:          Sun, 08 Sep 2019 12:54:12 -0400
Labels:              app=red-offer
                    pod-template-hash=2272044612
Annotations:         openshift.io/scc=restricted
Status:              Pending
IP:
Controlled By:       ReplicaSet/red-offer-66c6488b56
Containers:
  red-offer:
    Container ID:
    Image:         ibmicpcoc/offer:latest
    Image ID:
    Port:          <none>
    Host Port:     <none>
    Command:
      node
      app.js
    State:          Waiting
      Reason:       ContainerCreating
    Ready:          False
    Restart Count:  0
    Requests:
      cpu:          50m
      memory:       50Mi
    Environment:
      APP_NAMESPACE:  red (v1:metadata.namespace)
      APP_NAME:       red-offer-66c6488b56-m9wfr (v1:metadata.name)
      COLLECTOR_CONFIG: <set to the key 'COLLECTOR_CONFIG' of config map 'red-collector-config'> Optional:
false
      INSTRUCTOR_CONFIG: <set to the key 'INSTRUCTOR_CONFIG' of config map 'red-collector-config'> Optional:
false
    Mounts:
      /data from offer-data (rw)
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-dxnzt (ro)
    Conditions:
      Type          Status
```

```

    Initialized      True
    Ready            False
    ContainersReady  False
    PodScheduled     True
  Volumes:
    offer-data:
      Type:          PersistentVolumeClaim (a reference to a PersistentVolumeClaim in the same namespace)
      ClaimName:     red-offer
      ReadOnly:      false
    default-token-dxnzt:
      Type:          Secret (a volume populated by a Secret)
      SecretName:    default-token-dxnzt
      Optional:      false
  QoS Class:         Burstable
  Node-Selectors:    node-role.kubernetes.io/compute=true
  Tolerations:       node.kubernetes.io/memory-pressure:NoSchedule
  Events:
    Type      Reason      Age   From
    ----      -
    Normal    Scheduled    17s   default-scheduler
    Warning   FailedMount  16s   kubelet, gfstst.169.62.225.201.nip.io
    Message:  MountVolume.SetUp failed for volume "red-pv" : mount failed: exit status 32
    Mounting command: systemd-run
    Mounting arguments: --description=Kubernetes transient mount for
/var/lib/origin/openshift.local.volumes/pods/48d45636-d259-11e9-8c57-0607f5770d4d/volumes/kubernetes.io~nfs/red-pv -
-scope -- mount -t nfs 169.62.225.199:/storage/red/pvc001 /var/lib/origin/openshift.local.volumes/pods/48d45636-
d259-11e9-8c57-0607f5770d4d/volumes/kubernetes.io~nfs/red-pv
    Output: Running scope as unit run-126180.scope.
    mount.nfs: mounting 169.62.225.199:/storage/red/pvc001 failed, reason given by server: No such file or directory
    Warning   FailedMount  16s   kubelet, gfstst.169.62.225.201.nip.io
    Message:  MountVolume.SetUp failed for volume "red-pv" : mount failed: exit status 32

```

Problem 2 discovered

Message from the describe indicates the PV mount failed. This is caused because the path does not exist.

Resolution 2

Two options exist to correct this issue:

Option 1

Change the the PV nfs path to a path that exists.

Option 2 Create the path on the NFS server.

Delete and redeploy all resources i.e. PV, PVC, Pod, etc.

All references to "team" or <team> should be replaced with your team name which is the same as your namespace.

Desired environment

Deployment of an application that uses persistent storage. The storage is implemented as dynamic storage.

Note: **This lab requires the student to resolve multiple issues**

Resources

- K8 yaml - [panda.yaml](#)
- Dockerfile - [Dockerfile](#)

Useful information

Item	Value
cpu:	50m
memory:	50Mi
image:	ibmicpcoc/panda:latest
ports	none
YAML	command: ["node", "app.js"]

Task description

A statefulset that begins with <team>-panda is failing.

Research the issue to determine what is causing the statefulset to fail.

Reiview the pod log to determine how long the application http server waits to be started.

Edit the panda.yaml file to correct the issue.

Validate if the statefulset deployed.

Verify the pod deployed. If not research why not.

Edit the panda.yaml file to correct any issues.

Validate if the statefulset and pod deployed.

Did resource type is created in the yaml?

Ensure to review and diagnois all resource types.

Diagnosis 1

Checking the pod information.

```
Command to describe the statefulset
oc describe statefulset red-panda
```

Example output

```
Name:                red-panda
Namespace:           red
CreationTimestamp:    Sun, 08 Sep 2019 14:17:00 -0400
Selector:            app=red-panda
Labels:              app=red-panda
Annotations:         <none>
Replicas:            1 desired | 0 total
Update Strategy:     RollingUpdate
```

```

Pods Status:          0 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:  app=red-panda
  Containers:
    red-panda:
      Image:          ibmicpcoc/offer:latest
      Port:           <none>
      Host Port:      <none>
      Command:
        node
        app.js
      Requests:
        cpu:          50m
        memory:       50Mi
      Environment:
        APP_NAMESPACE: (v1:metadata.namespace)
        APP_NAME:       (v1:metadata.name)
        COLLECTOR_CONFIG: <set to the key 'COLLECTOR_CONFIG' of config map 'red-collector-config'> Optional:
false
        INSTRUCTOR_CONFIG: <set to the key 'INSTRUCTOR_CONFIG' of config map 'red-collector-config'> Optional:
false
      Mounts:
        /data from panda-d (rw)
      Volumes: <none>
  Volume Claims:
    Name:          panda-data
    StorageClass:  rdb
    Labels:        <none>
    Annotations:   <none>
    Capacity:      1Mi
    Access Modes:  [ReadWriteOnce]
  Events:
    Type      Reason      Age      From      Message
    ----      -
    Warning   FailedCreate  24s (x16 over 1m)  statefulset-controller  create Pod red-panda-0 in StatefulSet red-
panda failed error: Pod "red-panda-0" is invalid: spec.containers[0].volumeMounts[0].name: Not found: "panda-d"

```

Problem 1 discovered

The panda-d volumeMount does not exist.

Resolution 1

Ensure the parameters volumeClaimTemplate.metadata.name match the stateful set spec.template.spec.container.volumeMounts.name

Diagnosis 2

Checking the pod information.

```

Command to describe the statefulset
oc describe statefulset red-panda

```

```

Example output
Name:          red-panda-0

```

```

Namespace:      red
Priority:        0
PriorityClassName: <none>
Node:           <none>
Labels:         app=red-panda
                controller-revision-hash=red-panda-89c55dc87
                statefulset.kubernetes.io/pod-name=red-panda-0
Annotations:    openshift.io/scc=restricted
Status:         Pending
IP:
Controlled By:  StatefulSet/red-panda
Containers:
  red-panda:
    Image:        ibmicpcoc/offer:latest
    Port:         <none>
    Host Port:    <none>
    Command:
      node
      app.js
    Requests:
      cpu:        50m
      memory:     50Mi
    Environment:
      APP_NAMESPACE:  red (v1:metadata.namespace)
      APP_NAME:       red-panda-0 (v1:metadata.name)
      COLLECTOR_CONFIG: <set to the key 'COLLECTOR_CONFIG' of config map 'red-collector-config'> Optional:
false
      INSTRUCTOR_CONFIG: <set to the key 'INSTRUCTOR_CONFIG' of config map 'red-collector-config'> Optional:
false
    Mounts:
      /data from panda-data (rw)
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-dxnzt (ro)
    Conditions:
      Type           Status
      PodScheduled   False
    Volumes:
      panda-data:
        Type:          PersistentVolumeClaim (a reference to a PersistentVolumeClaim in the same namespace)
        ClaimName:     panda-data-red-panda-0
        ReadOnly:      false
      default-token-dxnzt:
        Type:          Secret (a volume populated by a Secret)
        SecretName:    default-token-dxnzt
        Optional:      false
    QoS Class:       Burstable
    Node-Selectors:  node-role.kubernetes.io/compute=true
    Tolerations:     node.kubernetes.io/memory-pressure:NoSchedule
    Events:
      Type          Reason          Age          From          Message
      ----          -
      Warning       FailedScheduling  52s (x25 over 1m)  default-scheduler  pod has unbound PersistentVolumeClaims
(repeated 3 times)

```


Command to get PVCs

```
oc get persistentvolumeclaims
```

Example output:

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES
STORAGECLASS	AGE			
panda-data-red-panda-0	Pending			rdb
2h				
red-panda-data-red-panda-0	Bound	pvc-9d13c3eb-d263-11e9-8c57-0607f5770d4d	1Mi	RWO
managed-nfs-storage	2h			

Command to describe the PVC

```
oc describe pvc panda-data-red-panda-0
```

Example output:

```
Name:          panda-data-red-panda-0
Namespace:     red
StorageClass:  rdb
Status:        Pending
Volume:
Labels:        app=red-panda
Annotations:   <none>
Finalizers:    [kubernetes.io/pvc-protection]
Capacity:
Access Modes:
Events:
  Type    Reason             Age          From                    Message
  ----    -
Warning  ProvisioningFailed  2m (x664 over 2h)  persistentvolume-controller  storageclass.storage.k8s.io "rdb"
not found
```

Problem 2 discovered

The pod has a unbound PCV. Getting the existing PVCs shows there is a Pending status. Describe the status of the pending PVC. The describe output shows the storage class 'rdb' does not exist.

Resolution 2

Determine the available storage classes and redefine the storage definition using a valid storage class.

Command to determine the available storage classes

```
oc get storageclass
```

Example output:

NAME	PROVISIONER	AGE
glusterfs-storage	kubernetes.io/glusterfs	10d

glusterfs-storage-block	gluster.org/glusterblock	10d
managed-nfs-storage	myokd/nfs	10d

All references to "team" or <team> should be replaced with your team name which is the same as your namespace.

Desired environment

Deploy a pod that is accessible external to the cluster via a route. The desired route name is defined as a environment variable. This environment variable does not create the route but defines what route must be defined.

Resources

- K8 yaml - [quake.yaml](#)
- Dockerfile - [Dockerfile](#)

Useful information

Item	Value
cpu:	50m
memory:	50Mi
image:	ibmicpcoc/quake:latest
ports	none
YAML	command: ["node", "app.js"]

Route Parameter	Value
metadata.labels.app	<team>-quake
metadata.name	hot-dog
metadata.namespace	<team>
spec.host	must be determined by student
spec.port.targetPort	<team>-quake
spec.to.kind	Service
spec.to.name	<team>-quake
spec.to.weight	100
spec.wildcardPolicy	None

Task description
A pod that begins with <team>-quake is Back-off restarting.
Research the issue to determine what is causing the pod to restart frequently.
Reiview the pod log to aid in determining what is causing the issue.

Task description

Edit the quake.yaml file to correct the issue.

Verify the deployment successfully deployed.

Define the missing route.

Diagnosis

Checking the running pod for application information.

Command to view pods

```
oc get pods
```

Example output

NAME	READY	STATUS	RESTARTS	AGE
red-quake-d5f9cb9bb-fmw75	1/1	Running	0	4s

Command to describe pods

```
oc describe po red-quake-d5f9cb9bb-fmw75
```

Example output

```
Name: red-quake-d5f9cb9bb-fmw75
Namespace: red
Priority: 0
PriorityClassName: <none>
Node: gfstst.169.62.225.207.nip.io/169.62.225.207
Start Time: Sun, 08 Sep 2019 17:16:57 -0400
Labels: app=red-quake
        pod-template-hash=819576566
Annotations: openshift.io/scc=restricted
Status: Running
IP: 10.130.0.228
Controlled By: ReplicaSet/red-quake-d5f9cb9bb
Containers:
  red-quake:
    Container ID: docker://00f351acce5c580fefba540e76291e68e2adaf76b8a7d503ed2ccb5ff41124f
    Image: ibmicpcoc/quake:v2
    Image ID: docker-
    pullable://docker.io/ibmicpcoc/quake@sha256:4412f897746e13d7941ca6ba4a2e5a15769de47e5c7970dcc73adb3efc608545
    Port: 4100/TCP
    Host Port: 0/TCP
    State: Terminated
      Reason: Error
      Exit Code: 1
      Started: Sun, 08 Sep 2019 17:17:03 -0400
      Finished: Sun, 08 Sep 2019 17:17:04 -0400
    Last State: Terminated
      Reason: Error
      Exit Code: 1
      Started: Sun, 08 Sep 2019 17:17:00 -0400
      Finished: Sun, 08 Sep 2019 17:17:01 -0400
    Ready: False
```

```

Restart Count: 1
Requests:
  cpu:      50m
  memory:   50Mi
Environment:
  APP_NAMESPACE:      red (v1:metadata.namespace)
  APP_NAME:           red-quake-d5f9cb9bb-fmw75 (v1:metadata.name)
  COLLECTOR_CONFIG:   <set to the key 'COLLECTOR_CONFIG' of config map 'red-collector-config'> Optional:
false
  INSTRUCTOR_CONFIG:  <set to the key 'INSTRUCTOR_CONFIG' of config map 'red-collector-config'> Optional:
false
  ROUTE:              hotdog-red.gfstst.169.62.225.197.nip.io
Mounts:
  /var/run/secrets/kubernetes.io/serviceaccount from default-token-dxnzt (ro)
Conditions:
  Type           Status
  Initialized     True
  Ready           False
  ContainersReady False
  PodScheduled    True
Volumes:
  default-token-dxnzt:
    Type:          Secret (a volume populated by a Secret)
    SecretName:     default-token-dxnzt
    Optional:       false
  QoS Class:       Burstable
  Node-Selectors:  node-role.kubernetes.io/compute=true
  Tolerations:     node.kubernetes.io/memory-pressure:NoSchedule
Events:
  Type    Reason      Age           From              Message
  ----    -
  Normal   Scheduled    18s           default-scheduler  Successfully assigned red/red-
quake-d5f9cb9bb-fmw75 to gfstst.169.62.225.207.nip.io
  Normal   Pulling      13s (x2 over 16s)  kubelet, gfstst.169.62.225.207.nip.io  pulling image
"ibmicpcoc/quake:v2"
  Normal   Pulled       12s (x2 over 16s)  kubelet, gfstst.169.62.225.207.nip.io  Successfully pulled image
"ibmicpcoc/quake:v2"
  Normal   Created      12s (x2 over 16s)  kubelet, gfstst.169.62.225.207.nip.io  Created container
  Normal   Started      12s (x2 over 15s)  kubelet, gfstst.169.62.225.207.nip.io  Started container
  Warning   BackOff      10s              kubelet, gfstst.169.62.225.207.nip.io  Back-off restarting failed
container

Command to view logs of pod
oc logs red-quake-d5f9cb9bb-fmw75

Example output:
9/8/2019, 9:22:52 PM :: quak001i - Application random key: 63657248-92fd-434e-b31c-f610b279f8f8
9/8/2019, 9:22:52 PM :: quak003i - Environment APP_NAMESPACE: red
9/8/2019, 9:22:52 PM :: quak004i - Environment APP_NAME: Using random key = red-quake-d5f9cb9bb-fmw75
9/8/2019, 9:22:52 PM :: quak013i - Environment COLLECTOR_CONFIG: http://red-student-ui
9/8/2019, 9:22:52 PM :: quak014i - Environment INSTRUCTOR_CONFIG: http://dashboard.default
9/8/2019, 9:22:52 PM :: quak014i - Environment ROUTE: hotdog-red.gfstst.169.62.225.197.nip.io
9/8/2019, 9:22:52 PM :: jazz007i - Quake Server started, port: 4400

```

```
9/8/2019, 9:22:52 PM :: quak011i - Initial request to route
9/8/2019, 9:22:52 PM :: quak012e - Error getting to Route: http://hotdog-red.gfstst.169.62.225.197.nip.io
message: null
```

Problem discovered

The pod logs show error message labled with id quak012e. The route for the pod is not defined.

Resolution

Define the pod route with the provided information.

All references to "team" or <team> should be replaced with your team name which is the same as your namespace.

Resources

- K8 yaml - [rainey.yaml](#)
- Dockerfile - [Dockerfile](#)

Useful information

Item	Value
cpu:	50m
memory:	50Mi
image:	ibmicpcoc/rainey:latest
ports	none
YAML	command: ["node", "app.js"]
Misc	Application waits

Task description
A pod that begins with <team>-igloo is frequently restarting.
Research the issue to determine what is causing the pod to restart frequently.
Reiview the pod log to determine how long the application http server waits to be started.
Edit the igloo.yaml file to correct the issue.
Verify the deployment successfully deployed.
Get the NodePort for the red-igloo service.
Get the IP address for the master node.
Using the above NodePort and the master ip address access the url: http://:

Provide a hint.

Diagnosis

Checking the running pod for application information.

```
Command to view the pod logs
oc logs red-igloo-5dd5b6c7b8-jqdvx

Example output
9/4/2019, 1:54:54 AM :: igloo900i - Waiting 10 seconds to start HTTP server
```

Problem discovered

Describe the problem.

Resolution

Describe the resolution.

All references to "team" or <team> should be replaced with your team name which is the same as your namespace.

Resources

- K8 yaml - [salty.yaml](#)
- Dockerfile - [Dockerfile](#)

Useful information

Item	Value
cpu:	50m
memory:	50Mi
image:	ibmicpcoc/salty:latest
ports	none
YAML	command: ["node", "app.js"]
Misc	Application waits

Task description
A pod that begins with <team>-igloo is frequently restarting.
Research the issue to determine what is causing the pod to restart frequently.
Reiview the pod log to determine how long the application http server waits to be started.
Edit the igloo.yaml file to correct the issue.
Verify the deployment successfully deployed.
Get the NodePort for the red-igloo service.
Get the IP address for the master node.
Using the above NodePort and the master ip address access the url: http://:

Provide a hint.

Diagnosis

Checking the running pod for application information.

Command to view the pod logs

```
oc logs red-igloo-5dd5b6c7b8-jqdv
```

Example output

```
9/4/2019, 1:54:54 AM :: igloo900i - Waiting 10 seconds to start HTTP server
```

Problem discovered

Describe the problem.

Resolution

Describe the resolution.
